

Tutorial on the MPL interface to MPI

Victor Eijkhout `eijkhout@tacc.utexas.edu`
TACC APP institute MPI training 2021

Justification

While the C API to MPI is usable from C++, it feels very unidiomatic for the language. Message Passing Layer (MPL) is a modern C++11 interface to MPI. It is both idiomatic and elegant, simplifying many calling sequences.

Part I

Basics

1. Environment

For doing the exercises:

```
module load mpl
```

which defines TACC_MPL_INC, TACC_MPL_DIR

2. Header file

To compile MPL programs, add a line

```
1  #include <mpl/mpl.hpp>
```

to your file.

3. Init, finalize

There is no initialization or finalize call.

*MPL implementation note: Initialization is done at the first
mpl::environment method call, such as comm_world.*

4. World communicator

The naive way of declaring a communicator would be:

```
1  // commrank.cxx
2  mpl::communicator comm_world =
3  mpl::environment::comm_world();
```

calling the predefined environment method `comm_world`.

However, if the variable will always correspond to the world communicator, it better to make it `const` and declare it to be a reference:

```
1  const mpl::communicator &comm_world =
2  mpl::environment::comm_world();
```

5. Processor name

The `processor_name` call is an environment method returning a `std::string`:

```
1 std::string mpl::environment::processor_name ();
```


6. Rank and size

The rank of a process (by `mpl::communicator::rank`) and the size of a communicator (by `mpl::communicator::size`) are both methods of the `communicator` class:

```
1  const mpl::communicator &comm_world =  
2      mpl::environment::comm_world();  
3  int procid = comm_world.rank();  
4  int nprocs = comm_world.size();
```

7. Reference: MPI_Comm_size

```
int mpl::communicator::size ( ) const
```

8. Reference: MPI_Comm_rank

```
int mpl::communicator::rank ( ) const
```

9. Timing

The timing routines `wtime` and `wtick` and `wtime_is_global` are environment met

```
1 double mpl::environment::wtime ();  
2 double mpl::environment::wtick ();  
3 bool mpl::environment::wtime_is_global ();
```

10. Predefined communicators

The `environment` namespace has the equivalents of `MPI_COMM_WORLD` and `MPI_COMM_SELF`:

```
1  const communicator& mpl::environment::comm_world();  
2  const communicator& mpl::environment::comm_self();
```

There doesn't seem to be an equivalent of `MPI_COMM_NULL`.

11. Communicator copying

The communicator class has its copy operator deleted; however, copy initialization exists:

```
1 // commcompare.cxx
2 const mpl::communicator &comm =
3     mpl::environment::comm_world();
4 cout << "same: " << boolalpha << (comm==comm) << endl;
5
6 mpl::communicator copy =
7     mpl::environment::comm_world();
8 cout << "copy: " << boolalpha << (comm==copy) << endl;
9
10 mpl::communicator init = comm;
11 cout << "init: " << boolalpha << (init==comm) << endl;
```

(This outputs true/false/false respectively.)

MPL implementation note: The copy initializer performs an

MPI_Comm_dup.

12. Communicator duplication

Communicators can be duplicated but only during initialization. Copy assignment has been deleted. Thus:

```
1 // LEGAL:
2 mpl::communicator init = comm;
3 // WRONG:
4 mpl::communicator init;
5 init = comm;
```

13. Communicator passing

Pass communicators by reference to avoid communicator duplication:

```
1 // commpass.cxx
2 // BAD! this does a MPI_Comm_dup.
3 void comm_val( const mpl::communicator comm );
4
5 // correct!
6 void comm_ref( const mpl::communicator &comm );
```

Part II

Collectives

Introduction

Collectives have many polymorphic variants, for instance for 'in place', and for
handling.

Operators are handled through functors.

14. Scalar buffers

Buffer type handling is done through polymorphism and templating: no explicit indication of types.

Scalars are handled as such:

```
1 float x,y;  
2 comm.bcast( 0,x ); // note: root first  
3 comm.allreduce( mpl::plus<float>(), x,y ); // op first
```

where the reduction function needs to be compatible with the type of the buffer

15. Vector buffers

If your buffer is a `std::vector` you need to take the `.data()` component of it:

```
1 vector<float> xx(2),yy(2);  
2 comm.allreduce( mpl::plus<float>(),  
3               xx.data(), yy.data(), mpl::contiguous_layout<float>(2) );
```

The `contiguous_layout` is a 'derived type'; this will be discussed in more detail elsewhere (see note 62 and later). For now, interpret it as a way of indicating count/type part of a buffer specification.

16. Iterator buffers

MPL point-to-point routines have a way of specifying the buffer(s) through a *begin* and *end* iterator.

```
1 // sendrange.cxx
2 vector<double> v(15);
3     comm_world.send(v.begin(), v.end(), 1); // send to rank 1
4     comm_world.recv(v.begin(), v.end(), 0); // receive from rank 0
```

Not available for collectives.

17. Reduction operator

The usual reduction operators are given as templated operators:

```
1    float
2        xrank = static_cast<float>( comm_world.rank() ),
3        xreduce;
4    // separate recv buffer
5    comm_world.allreduce(mpl::plus<float>(), xrank,xreduce);
6    // in place
7    comm_world.allreduce(mpl::plus<float>(), xrank);
```

Note the parentheses after the operator. Also note that the operator comes first, not last.

MPL implementation note: The reduction operator has to be compatible with $T(T,T)$

18. Reference: MPI_Allreduce

```
template<typename T , typename F >
void mpl::communicator::allreduce
    ( F,const T &, T & ) const;
    ( F,const T *, T *,
      const contiguous_layout< T > & ) const;
    ( F,T & ) const;
    ( F,T *, const contiguous_layout< T > & ) const;
F : reduction function
T : type
```

19. Reference: MPI_Reduce

```
void mpl::communicator::reduce
    // root, in place
    ( F f,int  root_rank,T &  sendrecvdata ) const
    ( F f,int  root_rank,T *  sendrecvdata,const contiguous_layout< T > & l ) const
    // non-root
    ( F f,int  root_rank,const T &  senddata ) const
    ( F f,int  root_rank,
        const T * senddata,const contiguous_layout< T > & l ) const
    // general
    ( F f,int  root_rank,const T &  senddata,T &  recvdata ) const
    ( F f,int  root_rank,
        const T *  senddata,T *  recvdata,const contiguous_layout< T > & l ) const
```


20. Broadcast

The broadcast call comes in two variants, with scalar argument and general la

```
1  template<typename T >
2  void mpl::communicator::bcast
3      ( int root_rank, T &data ) const;
4  void mpl::communicator::bcast
5      ( int root_rank, T *data, const layout< T > &l ) const;
```

Note that the root argument comes first.

21. Reference: MPI_Bcast

```
template<typename T >
void mpl::communicator::bcast
    ( int  root, T &  data ) const
    ( int  root, T *  data, const layout< T > &  l ) const
```

22. Gather scatter

Gathering (by `communicator::gather`) or scattering (by `communicator::scatter`) single scalar takes a scalar argument and a raw array:

```
1    vector<float> v;  
2    float x;  
3    comm_world.scatter(0, v.data(), x);
```

If more than a single scalar is gathered, or scattered into, it becomes necessary to specify a layout:

```
1    vector<float> vrecv(2), vsend(2*nprocs);  
2    mpl::contiguous_layout<float> twonums(2);  
3    comm_world.scatter  
4    (0, vsend.data(), twonums, vrecv.data(), twonums );
```

23. Reduce on non-root processes

There is a separate variant for non-root usage of rooted collectives:

```
1  // scangather.cxx
2  if (procno==0) {
3      comm_world.reduce
4          ( mpl::plus<int>(),0,
5            my_number_of_elements,total_number_of_elements );
6  } else {
7      comm_world.reduce
8          ( mpl::plus<int>(),0,my_number_of_elements );
9  }
```

24. Gather on non-root

Logically speaking, on every nonroot process, the gather call only has a send buffer. MPL supports this by having two variants that only specify the send c

```
1  if (procno==0) {
2      vector<int> size_buffer(nprocs);
3      comm_world.gather
4          (
5              0,my_number_of_elements,size_buffer.data()
6          );
7  } else {
8      /*
9       * If you are not the root, do versions with only send buffers
10      */
11      comm_world.gather
12          ( 0,my_number_of_elements );
```

25. Reference: MPI_Gather

```
void mpl::communicator::gather
( int  root_rank, const T & senddata ) const
( int  root_rank, const T & senddata, T *  recvdata ) const
( int  root_rank, const T * senddata, const layout< T > &  sendl
( int  root_rank, const T * senddata, const layout< T > &  sendl
                                T *  recvdata, const layout< T > &  recvl
```

26. Reduce in place

The in-place variant is activated by specifying only one instead of two buffer arguments.

```
1  float
2      xrank = static_cast<float>( comm_world.rank() ),
3      xreduce;
4  // separate recv buffer
5  comm_world.allreduce(mpl::plus<float>(), xrank,xreduce);
6  // in place
7  comm_world.allreduce(mpl::plus<float>(), xrank);
```

Reducing a buffer requires specification of a `contiguous_layout`:

```
1  // collectbuffer.cxx
2  float
3      xrank = static_cast<float>( comm_world.rank() );
4  vector<float> rank2p2p1{ 2*xrank,2*xrank+1 },reduce2p2p1{0,0};
5  mpl::contiguous_layout<float> two_floats(rank2p2p1.size());
6  comm_world.allreduce
7      (mpl::plus<float>(), rank2p2p1.data(),reduce2p2p1.data(),two_floats);
8  if ( iprint )
9      cout << "Got: " << reduce2p2p1.at(0) << ", "
10          << reduce2p2p1.at(1) << endl;
```

27. Layouts for gatherv

The size/displacement arrays for `MPI_Gatherv` / `MPI_Alltoallv` are handled through a `layouts` object, which is basically a vector of `layout` objects.

```
1      mpl::layouts<int> receive_layout;  
2      for ( int iproc=0, loc=0; iproc<nprocs; iproc++ ) {  
3          auto siz = size_buffer.at(iproc);  
4          receive_layout.push_back  
5              ( mpl::indexed_layout<int>( {{ siz, loc }} ) );  
6          loc += siz;  
7      }
```


28. Scan operations

As in the C/F interfaces, MPL interfaces to the scan routines have the same calling sequences as the 'Allreduce' routine.

Exercise 1 (scangather)

- Let each process compute a random value n_{local} , and allocate an array of length. Define

$$N = \sum n_{\text{local}}$$

- Fill the array with consecutive integers, so that all local arrays, laid end-to-end, contain the numbers $0 \cdots N - 1$. (See figure ??.)

Exercise 2 (scangather)

Take the code from exercise 1 and extend it to gather all local buffers onto rank zero. Since the local arrays are of differing lengths, this requires `MPI_Gatherv`.

How do you construct the lengths and displacements arrays?

29. Operators

Arithmetic: *plus, multiplies, max, min.*

Logic: *logical_and, logical_or, logical_xor.*

Bitwise: *bit_and, bit_or, bit_xor.*

30. User defined operators

A user-defined operator can be a templated class with an `operator()`. Example

```
1 // reduceuser.cxx
2 template<typename T>
3 class lcm {
4 public:
5     T operator()(T a, T b) {
6         T zero=T();
7         T t((a/gcd(a, b))*b);
8         if (t<zero)
9             return -t;
10        return t;
11    }
```

```
1     comm_world.reduce(lcm<int>(), 0, v, result);
```

31. Lambda reduction operators

You can also do the reduction by lambda:

```
1      comm_world.reduce
2      ( [] (int i,int j) -> int
3          { return i+j; },
4          0,data );
```

32. Nonblocking collectives

Nonblocking collectives have the same argument list as the corresponding blocking variant, except that instead of a `void` result, they return an `irequest`. (See 52)

```
1 // ireducescalar.cxx
2 float x{1.}, sum;
3 auto reduce_request =
4     comm_world.ireduce(mpl::plus<float>(), 0, x, sum);
5 reduce_request.wait();
6 if (comm_world.rank()==0) {
7     std::cout << "sum = " << sum << '\n';
8 }
```

Part III

Point-to-point communication

33. Blocking send and receive

MPL uses a default value for the tag, and it can deduce the type of the buffer.
Sending a scalar becomes:

```
1 // sendscalar.cxx
2 if (comm_world.rank()==0) {
3     double pi=3.14;
4     comm_world.send(pi, 1); // send to rank 1
5     cout << "sent: " << pi << '\n';
6 } else if (comm_world.rank()==1) {
7     double pi=0;
8     comm_world.recv(pi, 0); // receive from rank 0
9     cout << "got : " << pi << '\n';
10 }
```

(See also note 19.)

34. Sending arrays

MPL can send *static arrays* without further layout specification:

```
1 // sendarray.cxx
2 double v[2][2][2];
3     comm_world.send(v, 1); // send to rank 1
4     comm_world.recv(v, 0); // receive from rank 0
```

Sending vectors uses a general mechanism:

```
1 // sendbuffer.cxx
2 std::vector<double> v(8);
3 mpl::contiguous_layout<double> v_layout(v.size());
4     comm_world.send(v.data(), v_layout, 1); // send to rank 1
5     comm_world.recv(v.data(), v_layout, 0); // receive from rank 0
```

(See also note 20.)

35. Reference: MPI_Send

```
template<typename T >
void mpl::communicator::send
    ( const T scalar&,int dest,tag = tag(0) ) const
    ( const T *buffer,const layout< T > &,int dest,tag = tag(0) ) const
    ( iterT begin,iterT end,int dest,tag = tag(0) ) const
T : scalar type
begin : begin iterator
end : end iterator
```

36. Reference: MPI_Recv

```
template<typename T >
status mpl::communicator::recv
    ( T &,int,tag = tag(0) ) const inline
    ( T *,const layout< T > &,int,tag = tag(0) ) const
    ( iterT begin,iterT end,int source, tag t = tag(0) ) const
```

37. Tag types

Tag are *int* or an *enum* typ:

```
1      template<typename T >
2      tag_t (T t);
3      tag_t (int t);
4
```

Example:

```
1  // inttag.cxx
2  enum class tag : int { ping=1, pong=2 };
3  int pinger = 0, ponger = world.size()-1;
4  if (world.rank()==pinger) {
5      world.send(x, 1, tag::ping);
6      world.recv(x, 1, tag::pong);
7  } else if (world.rank()==ponger) {
8      world.recv(x, 0, tag::ping);
9      world.send(x, 0, tag::pong);
10 }
```

38. Message tag

MPL differs from other Application Programmer Interfaces (APIs) in its treatment of tags: a tag is not directly an integer, but an object of class `tag`.

```
1 // sendrecv.cxx
2     mpl::tag t0(0);
3     comm_world.sendrecv
4         ( mydata,sendto,t0,
5           leftdata,recvfrom,t0 );
```

The `tag` class has a couple of methods such as `mpl::tag::any()` (for the `MPI_ANY_TAG` wildcard in receive calls) and `mpl::tag::up()` (maximal tag, found the `MPI_TAG_UB` attribute).

39. Any source

The constant `mpl::any_source` equals `MPI_ANY_SOURCE` (by *constexpr*).

40. Send-recv call

The send-recv call in MPL has the same possibilities for specifying the send and receive buffer as the separate send and recv calls: scalar, layout, iterator. However, out of the nine conceivably possible routine signatures, only the versions are available where the send and receive buffer are specified the same way. Also, send and receive tag need to be specified; they do not have default values.

```
1  // sendrecv.cxx
2      mpl::tag t0(0);
3      comm_world.sendrecv
4          ( mydata,sendto,t0,
5            leftdata,recvfrom,t0 );
```


41. Status object

The `mpl::status_t` object is created by the receive (or wait) call:

```
1      mpl::contiguous_layout<double> target_layout(count);  
2      mpl::status_t recv_status =  
3          comm_world.recv(target.data(),target_layout, the_other);  
4      recv_count = recv_status.get_count<double>();
```

42. Status source querying

The `status` object can be queried:

```
1 int source = recv_status.source();
```

43. Receive count

The `get_count` function is a method of the status object. The argument type handled through templating:

```
1 // recvstatus.cxx
2 double pi=0;
3 auto s = comm_world.recv(pi, 0); // receive from rank 0
4 int c = s.get_count<double>();
5 std::cout << "got : " << c << " scalar(s): " << pi << '\n';
```

44. Requests from nonblocking calls

Nonblocking routines have an `irequest` as function result. Note: not a parameter passed by reference, as in the C interface. The various wait calls are methods of the `irequest` class.

```
1      double recv_data;  
2      mpl::irequest recv_request =  
3          comm_world.irecv( recv_data, sender );  
4      recv_request.wait();
```

You can not default-construct the request variable:

```
1  // DOES NOT COMPILE:  
2  mpl::irequest recv_request;  
3  recv_request = comm.irecv( ... );
```

This means that the normal sequence of first declaring, and then filling in, the request variable is not possible.

MPL implementation note: The wait call always returns a `status object`; not assigning it means that the destructor is called on it.

45. Request pools

Instead of an array of requests, use an `irequest_pool` object, which acts like a vector of requests, meaning that you can *push* onto it.

```
1 // irecvsource.cxx
2     mpl::irequest_pool recv_requests;
3     for (int p=0; p<nprocs-1; p++) {
4         recv_requests.push( comm_world.irecv( recv_buffer[p], p ) );
5     }
```

You can not declare a pool of a fixed size and assign elements.

46. Wait any

The `irequest_pool` class has methods `waitany`, `waitall`, `testany`, `testall`, `waitsome`, `testsome`.

The 'any' methods return a `std::pair<bool, size_t>`, with `false` meaning `index==MPI_UNDEFINED` meaning no more requests to be satisfied.

```
1      auto [success, index] = recv_requests.waitany();
2      if (success) {
3          auto recv_status = recv_requests.get_status(index);
```

Same for `testany`, then `false` means no requests test true.

Exercise 3 (setdiff)

Create two distributed arrays of positive integers. Take the set difference of the two: the first array needs to be transformed to remove from it those numbers that are in the second array.

How could you solve this with an `MPI_Allgather` call? Why is it not a good idea to do so? Solve this exercise instead with a circular bucket brigade algorithm.

47. Buffered send

Creating and attaching a buffer is done through `bsend_buffer` and a support routine `bsend_size` helps in calculating the buffer size:

```
1 // bufiring.cxx
2 vector<float> sbuf(BUFLEN), rbuf(BUFLEN);
3 int size{
4     ↪comm_world.bsend_size<float>(mpl::contiguous_layout<float>(BUFLEN))
5     mpl::bsend_buffer buff(size);
6     comm_world.bsend(sbuf.data(),mpl::contiguous_layout<float>(BUFLEN), nex
```

Constant: `mpl::bsend_overhead` is `constexpr`'d to the MPI constant `MPI_BSEND_OVERHEAD`.

48. Buffer attach and detach

There is a separate attach routine, but normally this is called by the constructor of the `bsend_buffer`. Likewise, the detach routine is called in the buffer destructor.

```
1 void mpl::environment::buffer_attach (void *buff, int size);  
2 std::pair< void *, int > mpl::environment::buffer_detach ();
```

49. Persistent requests

MPL returns a `prequest` from persistent 'init' routines, rather than an `irequest` (MPL note 52):

```
1  template<typename T >  
2  prequest send_init (const T &data, int dest, tag t=tag(0)) const;
```

Likewise, there is a `prequest_pool` instead of an `irequest_pool` (note 53).

Part IV

Derived Datatypes

50. Datatype handling

MPL mostly handles datatypes through subclasses of the *layout* class. Layout MPL routines are templated over the data type.

```
1 // sendlong.cxx
2 mpl::contiguous_layout<long long> v_layout(v.size());
3 comm.send(v.data(), v_layout, 1); // send to rank 1
```

Also works with complex of float and double.

The data types, where MPL can infer their internal representation, are enumeration types, C arrays of constant size and the template classes *std::array*, *std::pair* and *std::tuple* of the C++ Standard Template Library. The only limitation is, that the C array and the mentioned template classes hold data elements of types that can be sent or received by MPL.

51. Native MPI datatypes

Should you need the `MPI_Datatype` object contained in an MPL layout, there is access function `native_handle`.

52. Derived type handling

In MPL type creation routines are in the main namespace, templated over the datatypes.

```
1  // vector.cxx
2  vector<double>
3      source(stride*count);
4  if (procno==sender) {
5      mpl::strided_vector_layout<double>
6          newvectortype(count,1,stripe);
7      comm_world.send
8          (source.data(),newvectortype,the_other);
9  }
```

The commit call is part of the type creation, and freeing is done in the destruc

53. Contiguous type

The MPL interface makes extensive use of `contiguous_layout`, as it is the main way to declare a nonscalar buffer; see note 20.

54. Contiguous composing of types

Contiguous layouts can only use predefined types or other contiguous layouts their 'old' type. To make a contiguous type for other layouts, use `vector_layout`

```
1 // contiguous.cxx
2 mpl::contiguous_layout<int> type1(7);
3 mpl::vector_layout<int> type2(8, type1);
```

(Contrast this with `strided_vector_layout`; note 65.)

55. Vector type

MPL has the `strided_vector_layout` class as equivalent of the vector type:

```
1  // vector.cxx
2  vector<double>
3      source(stride*count);
4  if (procno==sender) {
5      mpl::strided_vector_layout<double>
6          newvectortype(count,1,stride);
7      comm_world.send
8          (source.data(),newvectortype,the_other);
9  }
```

(See note 64 for nonstrided vectors.)

56. Iterator buffers

MPL point-to-point routines have a way of specifying the buffer(s) through a *begin* and *end* iterator.

```
1 // sendrange.cxx
2 vector<double> v(15);
3     comm_world.send(v.begin(), v.end(), 1); // send to rank 1
4     comm_world.recv(v.begin(), v.end(), 0); // receive from rank 0
```

Not available for collectives.

57. Iterator layout

Noncontiguous iterable objects can be send with a `iterator_layout`:

```
1  std::list<int> v(20, 0);  
2  mpl::iterator_layout<int> l(v.begin(), v.end());  
3  comm_world.recv(&(*v.begin()), 1, 0);
```

(See also note 66.)

58. Subarray layout

The templated `subarray_layout` class is constructed from a vector of triplets of global size / subblock size / first coordinate.

```
1  mpl::subarray_layout<int>(  
2    { {ny, ny_1, ny_0}, {nx, nx_1, nx_0} }  
3  );
```

59. Indexed type

In MPL, the `indexed_layout` is based on a vector of 2-tuples denoting block length / block location.

```
1  // indexed.cxx
2  const int count = 5;
3  mpl::contiguous_layout<int>
4      fiveints(count);
5  mpl::indexed_layout<int>
6      indexed_where{ { {1,2}, {1,3}, {1,5}, {1,7}, {1,11} } };
7
8  if (procno==sender) {
9      comm_world.send( source_buffer.data(),indexed_where, receiver );
10 } else if (procno==receiver) {
11     auto recv_status =
12         comm_world.recv( target_buffer.data(),fiveints, sender );
13     int recv_count = recv_status.get_count<int>();
14     assert(recv_count==count);
15 }
```

60. Indexed block type

For the case where all block lengths are the same, use `indexed_block_layout`:

```
1 // indexedblock.cxx
2     mpl::indexed_block_layout<int>
3         indexed_where( 1, {2,3,5,7,11} );
4     comm_world.send( source_buffer.data(), indexed_where, receiver );
```

61. Struct type scalar

One could describe the MPI struct type as a collection of displacements, to be applied to any set of items that conforms to the specifications. An MPL `heterogeneous_layout` on the other hand, incorporates the actual data. Thus you could write

```
1 // structscalar.cxx
2 char c; double x; int i;
3 if (procno==sender) {
4     c = 'x'; x = 2.4; i = 37; }
5 mpl::heterogeneous_layout object( c,x,i );
6 if (procno==sender)
7     comm_world.send( mpl::absolute,object,receiver );
8 else if (procno==receiver)
9     comm_world.recv( mpl::absolute,object,sender );
```

Here, the `absolute` indicates the lack of an implicit buffer: the layout is absolute rather than a relative description.

62. Struct type general

More complicated data than scalars takes more work:

```
1  // struct.cxx
2  char c; vector<double> x(2); int i;
3  if (procno==sender) {
4      c = 'x'; x[0] = 2.7; x[1] = 1.5; i = 37; }
5  mpl::heterogeneous_layout object
6      ( c,
7        mpl::make_absolute(x.data(),mpl::vector_layout<double>(2)),
8        i );
9  if (procno==sender) {
10     comm_world.send( mpl::absolute,object,receiver );
11 } else if (procno==receiver) {
12     comm_world.recv( mpl::absolute,object,sender );
13 }
```

Note the `make_absolute` in addition to `absolute` mentioned above.

63. Extent resizing

Resizing a datatype does not give a new type, but does the resize 'in place':

```
1 void layout::resize(ssize_t lb, ssize_t extent);
```

Part V

Communicator manipulations

64. Communicator errhandler

MPL does not allow for access to the wrapped communicators. However, for `MPI_COMM_WORLD`, the routine `MPI_Comm_set_errhandler` can be called directly.

65. Communicator splitting

In MPL, splitting a communicator is done as one of the overloads of the communicator constructor;

```
1 // commsplit.cxx
2 // create sub communicator modulo 2
3 int color2 = procno % 2;
4 mpl::communicator comm2( mpl::communicator::split, comm_world, color2 )
5 auto procno2 = comm2.rank();
6
7 // create sub communicator modulo 4 recursively
8 int color4 = procno2 % 2;
9 mpl::communicator comm4( mpl::communicator::split, comm2, color4 );
10 auto procno4 = comm4.rank();
```

MPL implementation note: The `communicator::split` identifier is an object of class `communicator::split_tag`, itself is an otherwise empty subclass of `communicator`:

```
1 class split_tag {};
2 static constexpr split_tag split{};
```

66. Split by shared memory

Similar to ordinary communicator splitting 76: `communicator::split_shared`.

Part VI

Process topologies

67. Graph communicators

The constructor `dist_graph_communicator`

```
1  dist_graph_communicator  
2      (const communicator &old_comm, const source_set &ss,  
3          const dest_set &ds, bool reorder = true);
```

is a wrapper around `MPI_Dist_graph_create_adjacent`.

68. Graph communicator querying

Methods *indegree*, *outdegree* are wrappers around `MPI_Dist_graph_neighbors_c`. Sources and targets can be queried with *inneighbors* and *outneighbors*, which are wrappers around `MPI_Dist_graph_neighbors`.

Part VII

Other

69. Timing

The timing routines `wtime` and `wtick` and `wtime_is_global` are environment met

```
1 double mpl::environment::wtime ();  
2 double mpl::environment::wtick ();  
3 bool mpl::environment::wtime_is_global ();
```

70. Threading support

MPL always calls `MPI_Init_thread` requesting the highest level `MPI_THREAD_MULTIPLE`

```
1  enum mpl::threading_modes {
2      mpl::threading_modes::single = MPI_THREAD_SINGLE,
3      mpl::threading_modes::funneled = MPI_THREAD_FUNNELED,
4      mpl::threading_modes::serialized = MPI_THREAD_SERIALIZED,
5      mpl::threading_modes::multiple = MPI_THREAD_MULTIPLE
6  };
7  threading_modes mpl::environment::threading_mode ();
8  bool mpl::environment::is_thread_main ();
```

Missing from MPL

MPL is not a full MPI implementation

- File I/O
- One-sided communication
- Shared memory
- Process management