

# Parallel Computing for Science & Engineering

02/23/23 & ...

Instructors:

Lars Koesterke, TACC

Victor Eijkhout, TACC

# Scientific Computing Terminology

## Terms

## Definition

- NUMA
  - • Non Uniform Memory Access. In SMP systems with multiple CPUs access time to different parts of memory may vary.
- Affinity
  - • Propensity to maintain a process or thread on a hardware execution unit.
- SMP
  - • Symmetric Multi-Process(ing/or). Single OS system with shared memory.
- OpenMP
  - • Comment statement (F90) or #pragma (C/C++) that specifies parallel operations and control.
  - • The lexical extent that a directive controls.
- Runtime
  - • All code controlled by a directive— lexical extent + content of called routines.
  - • Code or a library within an executable that interacts with the operating system and can control code execution.

# OpenMP-- Overview

- Standard is ~25 years old. Mature language
- The “language” is easily comprehended.  
You can start simple and expand.
- Light Weight from System Perspective
- Very portable –GNU and vendor compilers.
- Spend time finding parallelism can be the most difficult part. The parallelism may be hidden.
- Writing Parallel OpenMP code examples is relatively easy.
- Developing parallel algorithms and/or parallelizing serial code is much harder.
- Expert level requires awareness of scoping and synchronization.
- Expansion into other performance relevant areas like: thread pinning and memory pinning

# OpenMP --- Shared Memory

- Shared Memory systems:
  - One Operating System
  - Instantiation of **ONE process**
  - **Threads are forked (created) from within your program.**
  - Multiple threads on multiple cores

# What is OpenMP (**Open Multi-Processing**)

- **De facto standard for Scientific Parallel Programming on Symmetric Multi-Processor (SMP) Systems.**
- **It is an API (Application Program Interface) for designing and executing parallel Fortran, C and C++ programs**
  - Based on threads, but
  - Higher-level than POSIX threads (Pthreads)  
(<http://www.llnl.gov/computing/tutorials/pthreads/#Abstract>)
- **Implemented by:**
  - **Pragmas/comments in code**
  - **Runtime Library (interface to OS and Program Environment)**
  - **Environment Variables**
- **Compiler option required to interpret/activate directives.**
- **<http://www.openmp.org/> has tutorials and description.**
- **Directed by OpenMP ARB (Architecture Review Board)**


# OpenMP History

- Primary OpenMP participants

AMD, Cray, Fujitsu, HP, IBM, Intel, NEC, PGI, Oracle, MS, TI, CAPS, NVIDIA  
ANL, LLNL, cOMPunity, EPCC, LANL, NASA, ORNL, RWTH, TACC

- OpenMP Fortran API, Version 1.0, 1997
- OpenMP C API, Version 1.0, 1998
- OpenMP 2.0 API for Fortran, 2000
- OpenMP 2.0 API for C/C++, 2002
- OpenMP 2.5 API for C/C++ & F90 2005
- OpenMP 3.0 Tasks May 2008
- OpenMP 3.1 July 2011
- OpenMP 4.0 **Affinity, Devices, Depend, SIMD** July 2013
- OpenMP 4.5
- OpenMP 5.0

# OpenMP History



**OpenMP 3.0: The World is still flat, no support for NUMA (yet)!**  
**OpenMP is hardware agnostic, it has no notion of data locality.**  
**The Affinity problem: How to maintain or improve the nearness of threads and their most frequently used data.**

**Or:**

**Where to run threads?**

**Where to place data?**

<http://terboven.wordpress.com/>

**Thread binding was added in OpenMP 4.0**

# Advantages/Disadvantages of OpenMP

- Pros
  - Shared Memory Parallelism is easier to learn.
  - Coarse-grained or fine-grained parallelism
  - Parallelization can be incremental
  - Widely available, portable
  - Converting serial code to OpenMP parallel can be easier than converting to MPI parallel.
  - Shared-Memory hardware is prevalent now.
    - Supercomputers **and** your desktop/laptop (and your phones)
    - GPUs (Graphics Cards), Multi-core CPUs
  - Complements MPI and enables full core utilization
- Cons
  - Scalability limited by memory architecture.
  - Available on Shared-Memory systems “only”.
  - Beware: “Upgrading” large serial code may be hard.



# OpenMP Parallel Directives

Supports parallelism by Directives in Fortran, C/C++,...

Unlike others that require base language changes and constructs

Unlike MPI which supports parallelism through communication lib.

OpenMP implementation through the compiler

Compiler optimizes OpenMP code

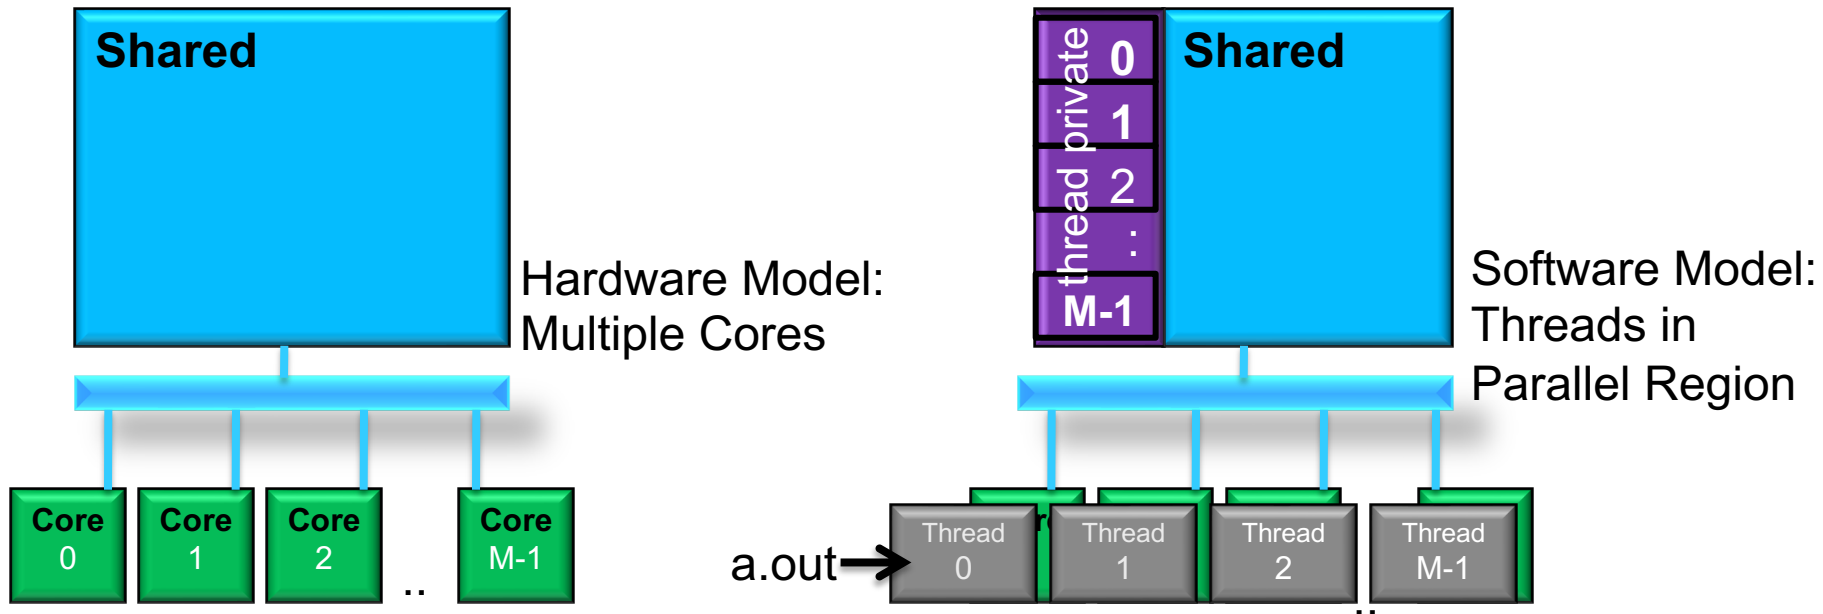
# Processes on a shared-memory System

- The OS starts a process
  - One instance of your computer program, the “a.out”
- Many processes may be executed on a single core through “time sharing” (time slicing).
  - The OS allows each process to run for awhile.
- The OS may run multiple processes concurrently on different cores.
- Security considerations
  - Independent processes have no direct communication (exchange of data) and are not able to read another process’s memory.
- Speed considerations
  - Time sharing among processes has a large overhead.

# OpenMP Threads

- Threads are instantiated (forked) in a program
- Threads run concurrently\*
- All threads (forked from the same process) can read the memory allocated to the process.
- Each thread is given some private memory only seen by the thread.
- \*When the # of threads forked exceeds the # of cores, time sharing (TS) will occur. Usually you would not do this. (But TS with user threads is less expensive than TS with processes).
- Implementation of threads differs from one OS to another.

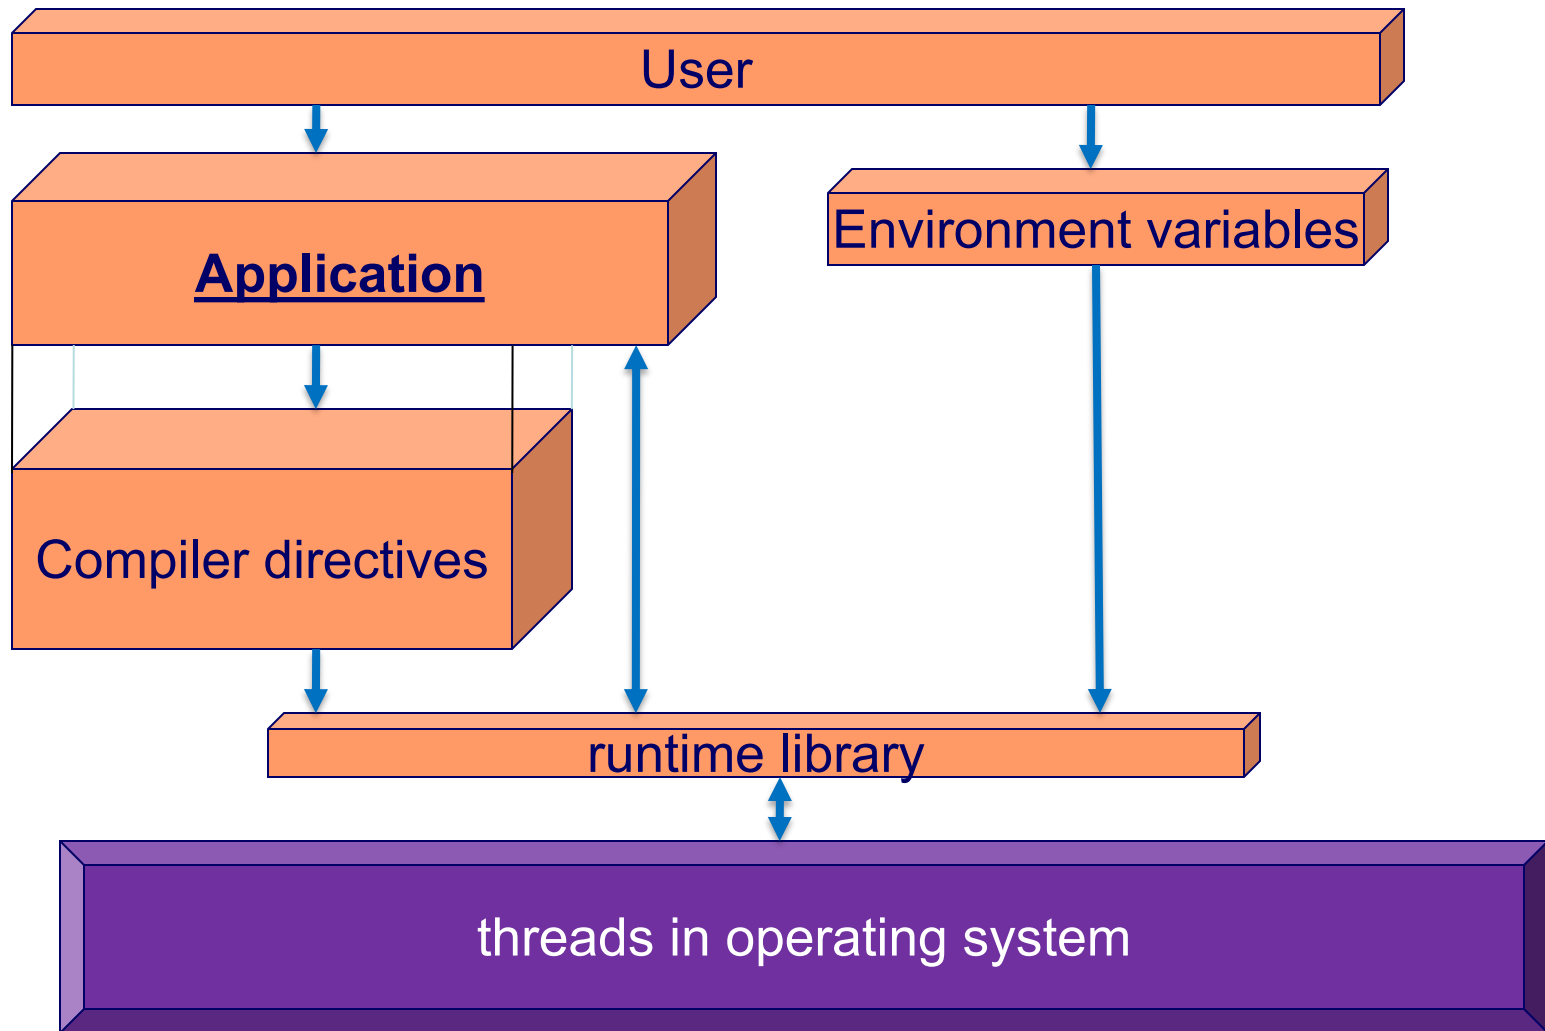
# Programming with OpenMP on Shared Memory Systems



M threads are usually mapped to M cores.

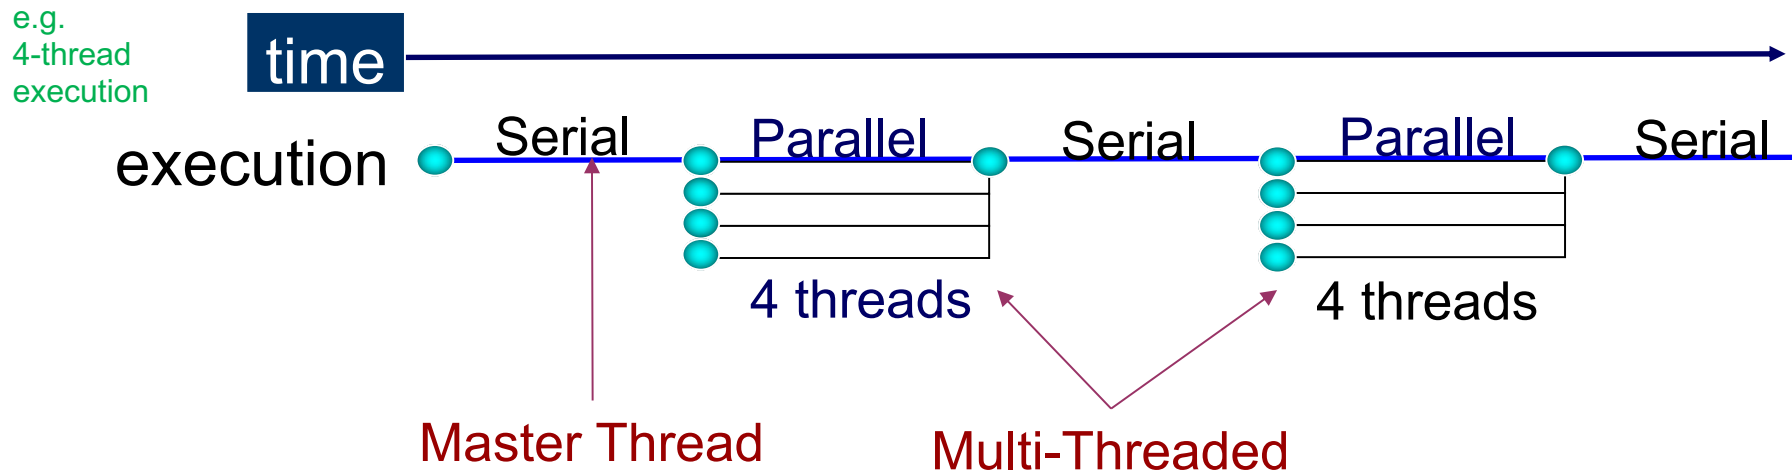
**Shared** = accessible by all threads  
**x** = private memory for thread x

# OpenMP Architecture



# OpenMP Fork-Join Parallelism

- **Programs begin as a single process:** master thread
- Master thread executes in serial mode until the parallel region construct is encountered
- Master thread **creates (forks)** a team of parallel **threads** that simultaneously execute tasks **in a parallel region**
- After executing the statements in the parallel region, team threads synchronize and terminate (**join**) but master continues



# Intermission

$$a = b + c$$



$$a(i) = b(i) + c(i)$$

```
loop (i) from 1 to n  
    a(i) = b(i) + c(i)  
end loop
```

# Executing in Parallel

How does this work?

```
loop (i) from 1 to n  
  a(i) = b(i) + c(i)  
end loop
```

You want to execute this with the help of your buddies  
What do you have to do?

# Executing in Parallel

How does this work?

```
loop (i) from 1 to n  
  a(i) = b(i) + c(i)  
end loop
```

You want to execute this with the help of your buddies  
What do you have to do?

1. Assemble a team
2. Divide up the work
3. Everybody works on their share of the loop
4. Wait for everybody to finish
5. Disassemble team

# Executing in Parallel

How does this work?

```
loop (i) from 1 to n  
  a(i) = b(i) + c(i)  
end loop
```

You want to execute this with the help of your buddies  
What do you have to do?

## ‘Plain english’

1. Assemble a team
2. Divide up the work
3. Everybody works
4. Wait for everybody
5. Disassemble team

## ‘OpenMP lingo’

Fork threads

Work sharing

Work in parallel

Barrier

Join threads

# OpenMP Syntax

- OpenMP Directives: **Sentinel**, **construct** and **clauses**

**#pragma omp construct ...** C

**!\$omp construct ...** F90

- Example for a loop

**#pragma omp parallel num\_threads(4)** C

**!\$omp parallel num\_threads(4)** F90

# Loop Example

```
...  
1. !$omp parallel  
1.  
2. !$omp do  
3. do i=1,n  
4.   a(i) = b(i)+c(i)  
5. end do  
6. !$omp end parallel
```

```
...  
#pragma omp parallel  
{  
#pragma omp for  
  for(i=0;i<n;i++){  
    a[i] = b[i]+c[i];  
  }  
}
```

**Let's identify the 5 steps:**

Fork threads, work-sharing, 'actual work', barrier, join threads

# Loop Example

```
...  
1. !$omp parallel  
1.  
2. !$omp do  
3. do i=1,n  
4.   a(i) = b(i)+c(i)  
5. end do  
6. !$omp end parallel
```

```
...  
#pragma omp parallel  
{  
#pragma omp for  
  for(i=0;i<n;i++){  
    a[i] = b[i]+c[i];  
  }  
}
```

**Let's identify the 5 steps:**

Fork threads, work-sharing, 'actual work', barrier, join threads

Forking/joining threads

Work-sharing with implied barrier at the end



## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?

How many work-sharing constructs?

How many barriers are there?

How many barriers do we need?

## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?  
How many barriers do we need

1 parallel region

2 work-sharing constructs

## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?  
How many barriers do we need

There is a barrier at the end of the parallel region

There is an implied barrier at the end of the every 'work-sharing' construct'

# 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?  
How many barriers do we need

There is a barrier at the end of the  
parallel region

There is an implied barrier at the end of the  
every 'work-sharing' construct'

Total number of barriers: 3

## 2<sup>nd</sup> Loop Example (modified)

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?

How many work-sharing constructs?

How many barriers are there?

**How many are necessary to ensure correct results?**

How many barriers do we need?



## 2<sup>nd</sup> Loop Example (modified)

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?

How many work-sharing constructs?

How many barriers are there?

**How many are necessary to ensure correct results?**

**How many barriers do we need?**

2 barriers are needed

First: to ensure correct results

Second: to join the threads



# End of Intermission

# Examples of Parallel Computing

- Concurrent execution of computational work (tasks).
  - Tasks execute independently
  - Variable Updates must be mutually exclusive
  - Synchronization through barriers

```
1  // We use loops for
   // repetitive tasks
2  for (i=0; i<N; i++){
3
4      a[i] = b[i] + c[i];
5  }
```

```
1  // We often update
   // variable(s)
2  for (i=0; i<N; i++){
3
4      sum = sum + b[i]*c[i];
5  }
```

Parallel directives go here ...



# OpenMP Constructs

OpenMP language  
“extensions”

parallel control  
structures

- governs flow of control in the program

**parallel**  
directive

parallel control  
worksharing

- distributes work among threads

**do**  
**for**  
**sections**  
**single**  
construct

control  
single task

- assigns work to a thread

**task**  
directive

data  
environment

- specifies variables as shared or private

**shared**  
**private**  
**reduction**  
clause

synchronization

- coordinates thread execution

**critical**  
**atomic**  
**barrier**  
directive

runtime  
environment

- runtime functions
- environment variables

**omp\_set\_num\_threads()**  
**omp\_get\_thread\_num()**  
**OMP\_NUM\_THREADS**  
**OMP\_SCHEDULE**

- scheduling  
**static, dynamic, guided**

# OpenMP Syntax

- OpenMP Directives: **Sentinel**, **construct** and **clauses**

```
#pragma omp construct [clause [,]clause]...]    C  
!$omp          construct [clause [,]clause]...]    F90
```

- Example

```
#pragma omp parallel num_threads(4)    C  
!$omp          parallel num_threads(4)    F90
```

- Function prototypes and types are in the file:

```
#include <omp.h>    C  
use omp_lib        F90
```

- Most OpenMP constructs apply to a “structured block”, that is, a block of one or more statements with one point of entry at the top and one point of exit at the bottom

# OpenMP Directives

- OpenMP **directives begin with special comments/pragmas** that open-aware compilers interpret. Directive **sentinels are**:

F90	<b>!\$OMP</b>
C/C++	<b># pragma omp</b>

Syntax: *sentinel* construct *clauses*      *defaults used when no clauses present*

## Fortran

```
!$OMP parallel
...
!$OMP end parallel
```

## C/C++

```
# pragma omp parallel
{ ... }
```

Fortran Parallel regions are enclosed by enclosing directives.

C/C++ Parallel regions are enclosed by curly brackets.

# Parallel Region

```
...  
1  !$omp parallel  
2      code statements  
3      call work(...)  
4  !$omp end parallel
```

```
...  
#pragma omp parallel  
{ code statements  
  work(...)  
}
```

Line 1      **Team of threads formed.**

Lines 2-3      This is the **parallel region**  
Each thread executes code block and  
subroutine call or function.

Line 4      No branching (in or out) in a parallel region.  
All threads **synchronize at end** of parallel  
region (implied barrier).

Replicated work or work-sharing?

# Parallel Region

```
...  
1  !$omp parallel  
2      code statements  
3      call work(...)  
4  !$omp end parallel
```

```
...  
#pragma omp parallel  
{  code statements  
  work(...)  
}
```

Line 1      **Team of threads formed.**

Lines 2-3    This is the **parallel region**  
Each thread executes code block and  
subroutine call or function.  
No branching (in or out) in a parallel region.

Line 4      All threads **synchronize at end** of parallel  
region (implied barrier).

In example above, user must explicitly create independent work (tasks) in the code block and routine (using thread id and total thread count).

# Parallel Region

```
...  
!$omp parallel  
  
    do i=1,n  
        call work(i)  
        a(i) = b(i)+c(i)  
    end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    for(i=0;i<n;i++){  
        work(i);  
        a[i] = b[i]+c[i];  
    }  
}
```

How often are the loop iterations executed?

Note: this code is not 100% correct (see discussion of private variables later)

# Parallel Region

```
...  
!$omp parallel  
  
    do i=1,n  
        call work(i)  
        a(i) = b(i)+c(i)  
    end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    for(i=0;i<n;i++){  
        work(i);  
        a[i] = b[i]+c[i];  
    }  
}
```

## Replicated work

In above example the do/for loop iterations **are not** split among the threads via a do/for work-sharing construct.

Note: this code is not 100% correct (see discussion of private variables later)

# Parallel Region with Worksharing Construct

```
...  
!$omp parallel  
  
!$omp do  
  do i=1,n  
    call work(i)  
    a(i) = b(i)+c(i)  
  end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
  #pragma omp for  
    for(i=0;i<n;i++){  
      work(i);  
      a[i] = b[i]+c[i];  
    }  
}
```

## Work-sharing

In above example the do/for loop iterations **are** split among the threads via the do/for worksharing constructs.



# Parallel Region & Work-Sharing

Use OpenMP directives to specify Parallel Region, Work-Sharing constructs, and Mutual Exclusion

`parallel`

`end parallel`

Use `parallel ... end parallel` for F90  
Use `parallel {...}` for C

`parallel do/for`  
`parallel sections`

*Code block*

`do / for`

`sections`

`single`

`master`

`critical`

`atomic`

**Each Thread Executes**

**Work Sharing**

Work Sharing

One Thread (Work sharing)

One Thread

One Thread at a time

One Thread at a time

A single worksharing construct (e.g. a `do/for`) may be combined on a parallel directive line.

# OpenMP Combined Directives

- Combined directives
  - `parallel do/for` and `parallel sections`
  - Same as parallel region containing only `do/for` or `sections` worksharing construct

```
!$omp parallel do  
  do i = 1, 100  
    a(i) = b(i)  
  end do
```

trip count required  
**no exit**  
cycle ok

```
#pragma omp parallel for  
for(i=0;i<100;i++){  
  a[i] = b[i];  
}
```

trip count required  
**no break**  
limited C++ throw.  
continue ok

# Work Sharing – do/for

## Worksharing (WS) constructs: do/for, sections, and single

- With Work-sharing: Threads execution their “share” of statements in a PARALLEL region.
- Do/for Work-sharing may require run-time work distribution and scheduling

1	!\$OMP PARALLEL DO	#pragma omp parallel for
2	do i=1,n	for(i=0;i<n;i++){
3	a(i)=b(i)+c(i)	a[i]=b[i]+c[i];
	enddo	}
5	!\$OMP END PARALLEL DO	

Line 1 Team of threads formed (parallel region).

Line 2-4 Loop iterations are split among threads.

**Implied barrier at “enddo” and “}”.**

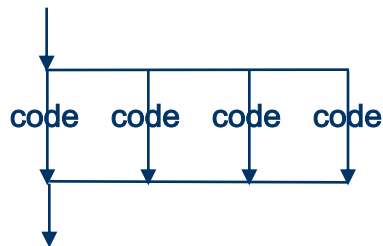
Line 5 (Optional) end of parallel loop.

- Each loop iteration must be independent of other iterations.

# Replicated and Work Share Constructs

- **Replicated:** Work blocks are executed by all threads.
- **Work Sharing:** Work is divided among threads.

```
PARALLEL  
  {code}  
END PARALLEL
```

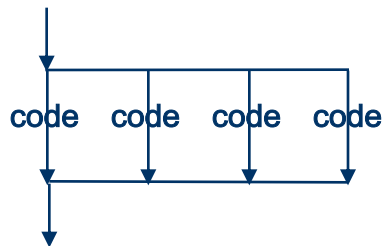


Replicated

# Replicated and Work Share Constructs

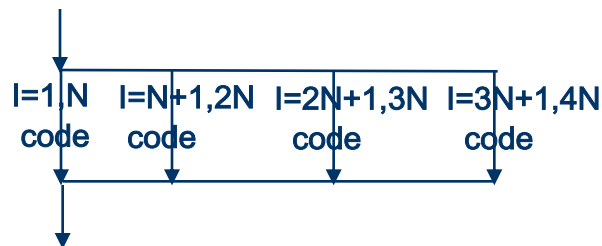
- **Replicated:** Work blocks are executed by all threads.
- **Work Sharing:** Work is divided among threads.

```
PARALLEL  
  {code}  
END PARALLEL
```



Replicated

```
PARALLEL DO  
  do I = 1,N*4  
    {code}  
  end do  
END PARALLEL DO
```



Work Sharing

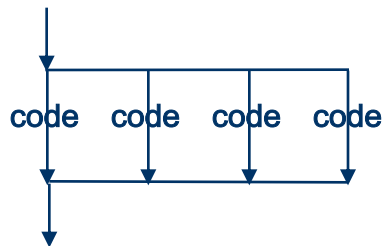
# Replicated and Work Share Constructs

- **Replicated:** Work blocks are executed by all threads.
- **Work Sharing:** Work is divided among threads.

```

PARALLEL
  { code }
END PARALLEL

```




Replicated

```

PARALLEL DO
  do I = 1,N*4
    {code}
  end do
END PARALLEL DO

```

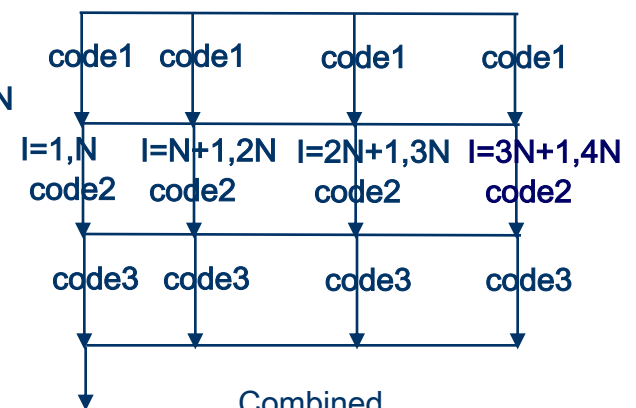


## Work Sharing

```

PARALLEL
    {code1}
DO
    do I = 1,N*4
        {code2}
    end do
        {code3}
END PARALLEL

```



## Combined

# OpenMP Worksharing Scheduling

Clause Syntax: **parallel do/for** **schedule**(*schedule-type* [,*chunk-size*])

## Schedule Type

### **schedule** (static[, chunk])

- Threads receive chunks of iterations in thread order, round-robin. (Divided “equally” if no chunk size.)
- Good if every iteration contains same amount of work
- May help keep parts of an array in a particular processor's cache— good between parallel do/for's.

### **schedule** (dynamic[, chunk])

- Thread receives chunks as it (the thread) becomes available for more work
- Default chunk size **may be** 1
- Good for load-balancing

# OpenMP Worksharing Scheduling

## `schedule (guided[, chunk])`

- Thread receives chunks as the thread becomes available for work
- Chunk size decreases logarithmically, until it reaches the chunk size specified (default is 1)
- Balances load and reduces number of requests for more work

## `schedule (runtime)`

- Schedule is determined at run-time by the `OMP_SCHEDULE` environment value.
- Useful for experimentation



# OpenMP Worksharing Scheduling

For example, loop with 100 iterations and 4 threads

- `schedule(static)`

Thread	0	1	2	3
Iteration	1-25	26-50	51-75	76-100

- `Schedule(static,10)`

Thread	0	1	2	3
Iteration	1-10, 41-50, 81-90	11-20, 51-60, 91-100	21-30, 61-70	31-40, 71-80

- `schedule(dynamic, 15)` (*one possible outcome*)

Thread	0	1	3	2	1	3	2
Iteration	1-15	16-30	31-45	46-60	61-75	76-90	90-100

- `schedule(guided, 8)` (*one possible outcome*)

Thread	0	1	2	3	3	2	3	1
Iteration	1-25	26-44	45-58	59-69	70-77	78-85	86-93	93-100

# Parallel – Worksharing - Schedule

- Combined directives
  - parallel do/for
  - Schedule clause added

```
!$omp parallel do schedule(static,8)  
do i = 1, 100  
  a(i) = b(i)  
end do
```

```
#pragma omp parallel for schedule(static,8)  
for(i=0;i<100;i++){  
  a[i] = b[i];  
}
```

How will the loop iteration be scheduled?  
Assume that the number of threads is 4

# OpenMP WorkSharing -- Sections

We will not discuss OpenMP sections in this class

- **SECTIONS**

- Blocks of code are split among threads - task parallel style
- A thread might execute more than one block or no blocks
- Implied barrier

**!\$OMP SECTIONS**

**!\$OMP SECTION**  
    **CALL TASK1()**

**!\$OMP SECTION**  
    **CALL TASK2()**

**!\$OMP SECTION**  
    **CALL TASK3()**

**!\$OMP END SECTIONS**

**#pragma omp sections**  
**{**

**#pragma omp section**  
    **{ TASK1( ); }**

**#pragma omp section**  
    **{ TASK2 ( ); }**

**#pragma omp section**  
    **{ TASK3 ( ); }**

**}**

# OpenMP Worksharing -- Single

- SINGLE (or MASTER)
  - Block of code is executed only once by a single thread (or the master thread)
  - Implied barrier (only SINGLE, not MASTER)

**!\$OMP single**

```
glob_count = glob_count + 1  
print *, glob_count
```

**!\$OMP end single**

**#pragma single**

```
{  
    glob_count++;  
    printf("%d\n", glob_count);
```

```
}
```

```
1 !$OMP PARALLEL DO
2   do i=1,n
3     a(i)=b(i)+c(i)
4   enddo
5 !$OMP END PARALLEL DO
```

```
#pragma omp parallel for
for(i=0;i<n;i++){
  a[i]=b[i]+c[i];
}
```

How many variables **i** are there?

Logical and 'in memory'