

# Parallel Computing for Science & Engineering

02/23/23 & ...

Instructors:

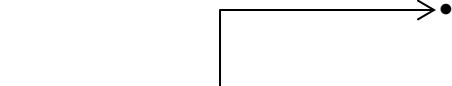
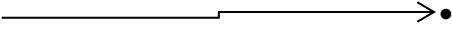

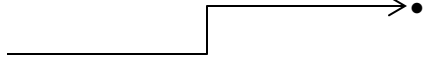

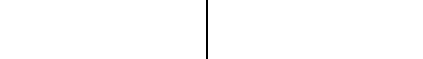
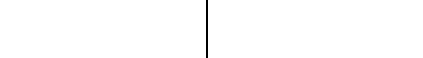

Lars Koesterke, TACC

Victor Eijkhout, TACC

# Scientific Computing Terminology

## Terms

## Definition

- NUMA • Non Uniform Memory Access. In SMP systems with multiple CPUs access time to different parts of memory may vary.
- Affinity • Propensity to maintain a process or thread on a hardware execution unit.
- SMP • Symmetric Multi-Process(ing/or). Single OS system with shared memory.
- OpenMP
  - Directive • Comment statement (F90) or #pragma (C/C++) that specifies parallel operations and control.
  - Construct • The lexical extent that a directive controls.
  - Region 
- Runtime
  - • All code controlled by a directive– lexical extent + content of called routines.
  - • Code or a library within an executable that interacts with the operating system and can control code execution.

# OpenMP-- Overview

- Standard is ~25 years old. Mature language
- The “language” is easily comprehended.  
You can start simple and expand.
- Light Weight from System Perspective
- Very portable –GNU and vendor compilers.
- Spend time finding parallelism can be the most difficult part. The parallelism may be hidden.
- Writing Parallel OpenMP code examples is relatively easy.
- Developing parallel algorithms and/or parallelizing serial code is much harder.
- Expert level requires awareness of scoping and synchronization.
- Expansion into other performance relevant areas like: thread pinning and memory pinning

# OpenMP --- Shared Memory

- Shared Memory systems:
  - One Operating System
  - Instantiation of **ONE process**
  - **Threads are forked (created) from within your program.**
  - Multiple threads on multiple cores

# What is OpenMP (**Open Multi-Processing**)

- **De facto standard for Scientific Parallel Programming on Symmetric Multi-Processor (SMP) Systems.**
- **It is an API (Application Program Interface) for designing and executing parallel Fortran, C and C++ programs**
  - Based on threads, but
  - Higher-level than POSIX threads (Pthreads)  
(<http://www.llnl.gov/computing/tutorials/pthreads/#Abstract>)
- **Implemented by:**
  - **Pragmas/comments in code**
  - **Runtime Library (interface to OS and Program Environment)**
  - **Environment Variables**
- **Compiler option required to interpret/activate directives.**
- **<http://www.openmp.org/> has tutorials and description.**
- **Directed by OpenMP ARB (Architecture Review Board)**


# OpenMP History

- Primary OpenMP participants

AMD, Cray, Fujitsu, HP, IBM, Intel, NEC, PGI, Oracle, MS, TI, CAPS, NVIDIA  
ANL, LLNL, cOMPunity, EPCC, LANL, NASA, ORNL, RWTH, TACC

- OpenMP Fortran API, Version 1.0, 1997
- OpenMP C API, Version 1.0, 1998
- OpenMP 2.0 API for Fortran, 2000
- OpenMP 2.0 API for C/C++, 2002
- OpenMP 2.5 API for C/C++ & F90 2005
- OpenMP 3.0 Tasks May 2008
- OpenMP 3.1 July 2011
- OpenMP 4.0 **Affinity, Devices, Depend, SIMD** July 2013
- OpenMP 4.5
- OpenMP 5.0

# OpenMP History



**OpenMP 3.0: The World is still flat, no support for NUMA (yet)!**  
**OpenMP is hardware agnostic, it has no notion of data locality.**  
**The Affinity problem: How to maintain or improve the nearness of threads and their most frequently used data.**

**Or:**

**Where to run threads?**

**Where to place data?**

<http://terboven.wordpress.com/>

**Thread binding was added in OpenMP 4.0**

# Advantages/Disadvantages of OpenMP

- Pros
  - Shared Memory Parallelism is easier to learn.
  - Coarse-grained or fine-grained parallelism
  - Parallelization can be incremental
  - Widely available, portable
  - Converting serial code to OpenMP parallel can be easier than converting to MPI parallel.
  - Shared-Memory hardware is prevalent now.
    - Supercomputers **and** your desktop/laptop (and your phones)
    - GPUs (Graphics Cards), Multi-core CPUs
  - Complements MPI and enables full core utilization
- Cons
  - Scalability limited by memory architecture.
  - Available on Shared-Memory systems “only”.
  - Beware: “Upgrading” large serial code may be hard.



# OpenMP Parallel Directives

Supports parallelism by Directives in Fortran, C/C++,...

Unlike others that require base language changes and constructs

Unlike MPI which supports parallelism through communication lib.

OpenMP implementation through the compiler

Compiler optimizes OpenMP code

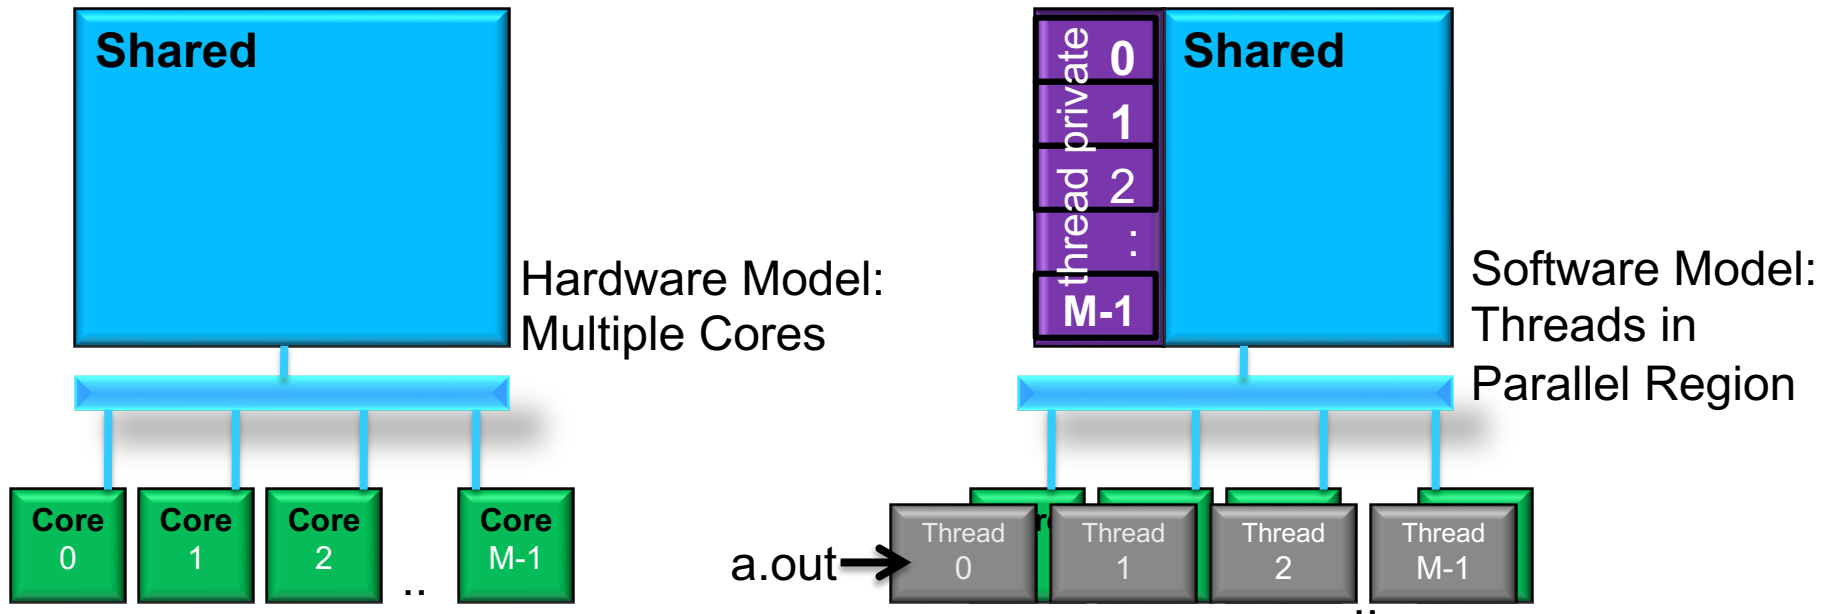
# Processes on a shared-memory System

- The OS starts a process
  - One instance of your computer program, the “a.out”
- Many processes may be executed on a single core through “time sharing” (time slicing).
  - The OS allows each process to run for awhile.
- The OS may run multiple processes concurrently on different cores.
- Security considerations
  - Independent processes have no direct communication (exchange of data) and are not able to read another process’s memory.
- Speed considerations
  - Time sharing among processes has a large overhead.

# OpenMP Threads

- Threads are instantiated (forked) in a program
- Threads run concurrently\*
- All threads (forked from the same process) can read the memory allocated to the process.
- Each thread is given some private memory only seen by the thread.
- \*When the # of threads forked exceeds the # of cores, time sharing (TS) will occur. Usually you would not do this. (But TS with user threads is less expensive than TS with processes).
- Implementation of threads differs from one OS to another.

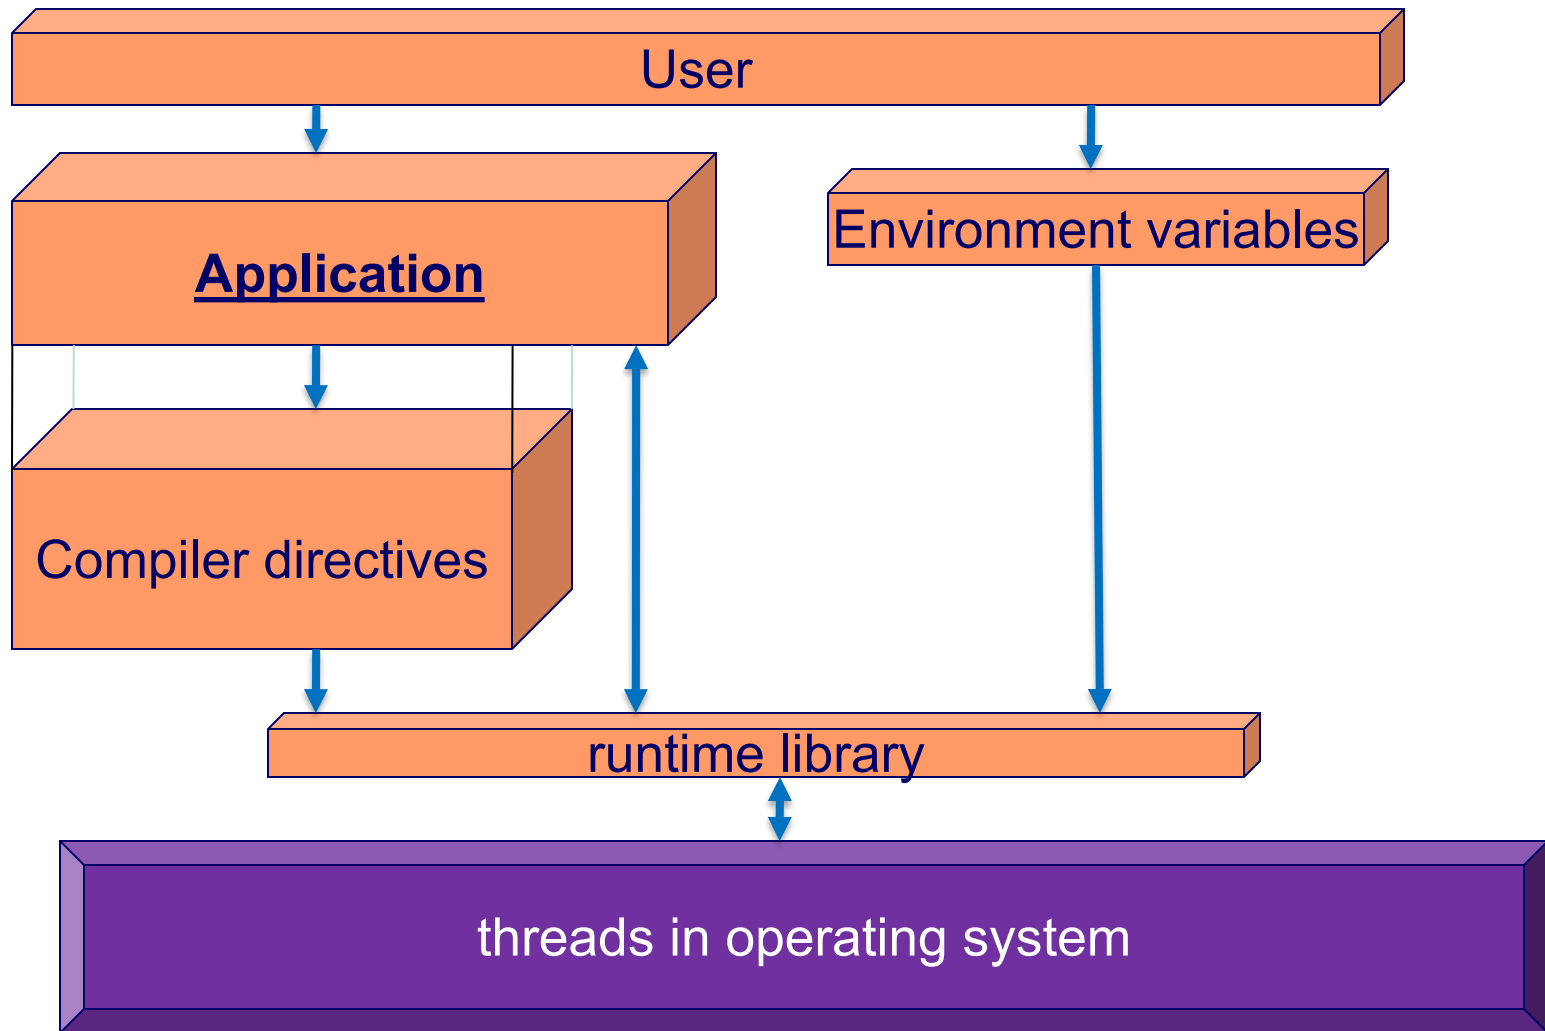
# Programming with OpenMP on Shared Memory Systems



M threads are usually mapped to M cores.

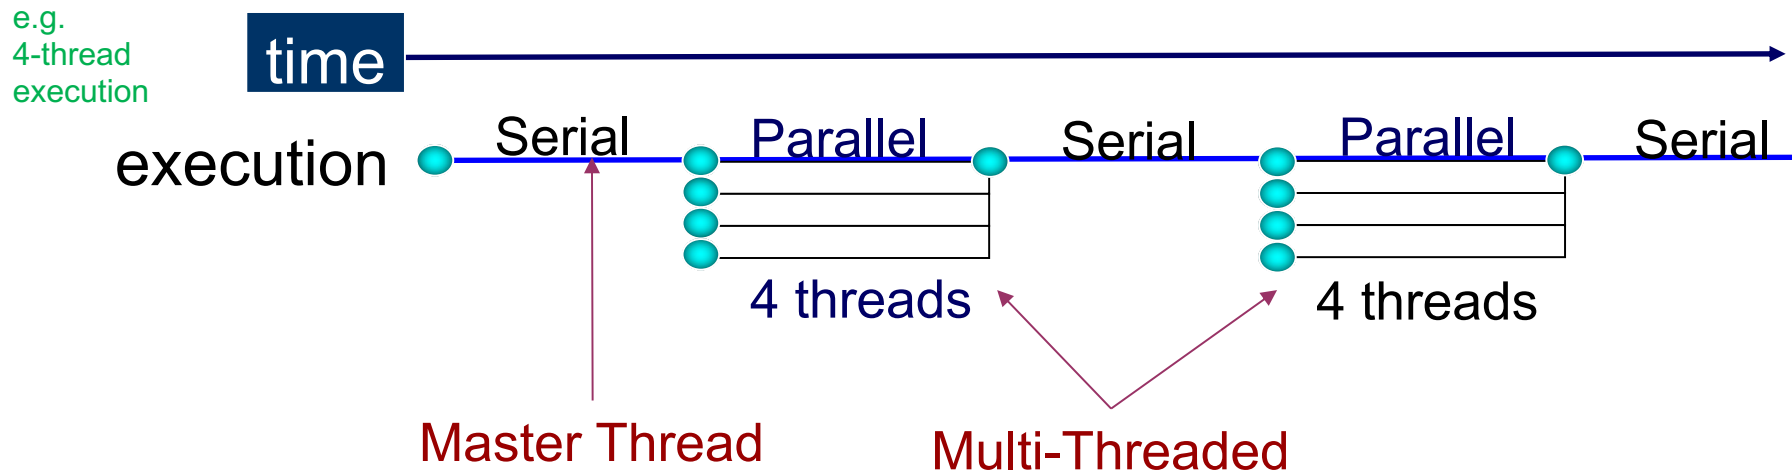
**Shared** = accessible by all threads  
**x** = private memory for thread x

# OpenMP Architecture



# OpenMP Fork-Join Parallelism

- **Programs begin as a single process:** master thread
- Master thread executes in serial mode until the parallel region construct is encountered
- Master thread **creates (forks)** a team of parallel **threads** that simultaneously execute tasks **in a parallel region**
- After executing the statements in the parallel region, team threads synchronize and terminate (**join**) but master continues



# Intermission

$$a = b + c$$



$$a(i) = b(i) + c(i)$$

```
loop (i) from 1 to n  
    a(i) = b(i) + c(i)  
end loop
```

# Executing in Parallel

How does this work?

```
loop (i) from 1 to n  
  a(i) = b(i) + c(i)  
end loop
```

You want to execute this with the help of your buddies  
What do you have to do?

# Executing in Parallel

How does this work?

```
loop (i) from 1 to n  
  a(i) = b(i) + c(i)  
end loop
```

You want to execute this with the help of your buddies  
What do you have to do?

1. Assemble a team
2. Divide up the work
3. Everybody works on their share of the loop
4. Wait for everybody to finish
5. Disassemble team

# Executing in Parallel

How does this work?

```
loop (i) from 1 to n
  a(i) = b(i) + c(i)
end loop
```

You want to execute this with the help of your buddies  
What do you have to do?

## 'Plain english'

1. Assemble a team
2. Divide up the work
3. Everybody works
4. Wait for everybody
5. Disassemble team

## 'OpenMP lingo'

Fork threads

Work sharing

Work in parallel

Barrier

Join threads

# OpenMP Syntax

- OpenMP Directives: **Sentinel**, **construct** and **clauses**

**#pragma omp construct ...** C

**!\$omp construct ...** F90

- Example for a loop

**#pragma omp parallel num\_threads(4)** C

**!\$omp parallel num\_threads(4)** F90

# Loop Example

```
...  
1. !$omp parallel  
1.  
2. !$omp do  
3. do i=1,n  
4.   a(i) = b(i)+c(i)  
5. end do  
6. !$omp end parallel
```

```
...  
#pragma omp parallel  
{  
#pragma omp for  
  for(i=0;i<n;i++){  
    a[i] = b[i]+c[i];  
  }  
}
```

**Let's identify the 5 steps:**

Fork threads, work-sharing, 'actual work', barrier, join threads

# Loop Example

```
...  
1. !$omp parallel  
1.  
2. !$omp do  
3. do i=1,n  
4.   a(i) = b(i)+c(i)  
5. end do  
6. !$omp end parallel
```

```
...  
#pragma omp parallel  
{  
#pragma omp for  
  for(i=0;i<n;i++){  
    a[i] = b[i]+c[i];  
  }  
}
```

**Let's identify the 5 steps:**

Fork threads, work-sharing, 'actual work', barrier, join threads

Forking/joining threads

Work-sharing with implied barrier at the end



## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?

How many work-sharing constructs?

How many barriers are there?

How many barriers do we need?

## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?  
How many barriers do we need

1 parallel region

2 work-sharing constructs

## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?  
How many barriers do we need

There is a barrier at the end of the  
parallel region

There is an implied barrier at the end of the  
every 'work-sharing' construct'

## 2<sup>nd</sup> Loop Example

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n/2  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n/2;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?  
How many barriers do we need

There is a barrier at the end of the  
parallel region

There is an implied barrier at the end of the  
every 'work-sharing' construct'

Total number of barriers: 3

## 2<sup>nd</sup> Loop Example (modified)

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
#pragma omp for  
for(i=0;i<n;i++){  
    a[i] = b[i]+c[i];  
}  
  
#pragma omp for  
for(i=0;i<n;i++){  
    a[i] = a[i]*d[i];  
}  
}
```

How many parallel regions?

How many work-sharing constructs?

How many barriers are there?

**How many are necessary to ensure correct results?**

How many barriers do we need?



## 2<sup>nd</sup> Loop Example (modified)

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    a(i) = b(i)+c(i)  
end do  
  
!$omp do  
do i=1,n  
    a(i) = a(i)*d(i)  
end do  
!$omp end parallel
```

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = b[i]+c[i];  
    }  
  
    #pragma omp for  
    for(i=0;i<n;i++){  
        a[i] = a[i]*d[i];  
    }  
}
```

How many parallel regions?  
How many work-sharing constructs?

How many barriers are there?

**How many are necessary to ensure correct results?**

**How many barriers do we need?**

2 barriers are needed

First: to ensure correct results  
Second: to join the threads



# End of Intermission